



# Generative AI Governance Strategy at Scale

Brian Scott | Principal Architect

Dan Neff | Principal Cloud Architect



Artwork by Leandro Alzate

# 41 Years of Innovation



# Created with Adobe Firefly



Adobe is one of the largest and most diversified software companies in the world



**30,000+**

Employees in  
35 countries



**41**

Years of  
revolutionizing industries



**\$19.41B**

FY2023  
Revenue

**85%**

Of our global  
workspaces are LEED-  
certified

**8,204**

Patents\*

**~\$95.5M**

Given to the  
community in 2021

# Who we are and what we've been doing

Brian  
Scott



Dan Neff



Principal Cloud Architects  
Adobe ⇒ Tech Services ⇒ Cloud Operations

Decades of operations experience (yay, Disney &  
Friendster!)

Lead tech onboarding for Dev/Eng tools, Open Source

Based on onboarding experience with finance, legal, security we were appointed ownership of company-wide AI intake *in addition to existing responsibilities*.

Created process that required no dedicated staff for first year of operation.  
(currently 1 full-time employee)

# Enterprise technology tension

Maximum Liberty

Solve customer needs



Best in breed tech solutions



POCs, prototypes,  
hackathons

Add-ons to existing standard

Maximum Responsibility

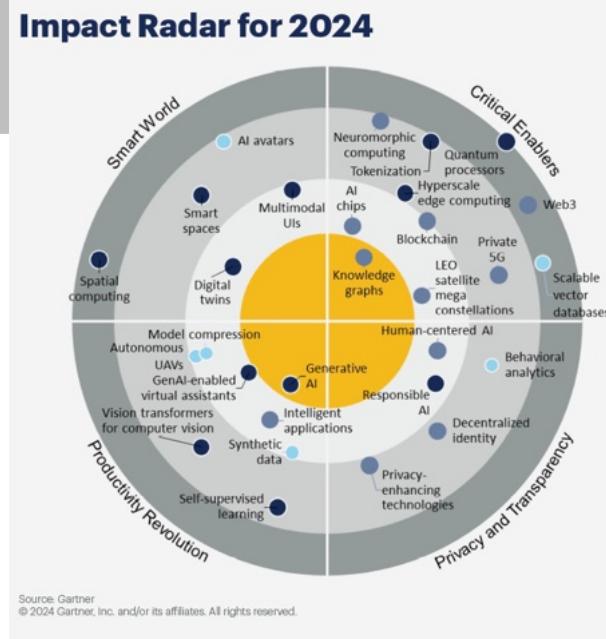
Compliance with security  
standards



Friendly contracts, licensing,  
EULA, TOUs



Aligns with data governance  
policies



# AI use case definition and value



Team [A] → ***The customer support engineering team***



with objective [B] → ***to improve customer quality of service***



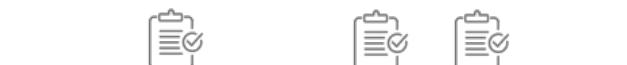
for audience [C] → ***for internal customer service reps***



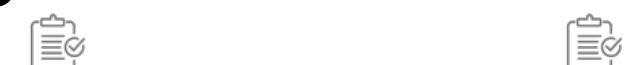
uses input data [D] → ***using published documentation and  
internal runbooks***



with tech [E] → ***with Azure OpenAI, LangChain, Pinecone***



to create output [F] → ***to find concrete, referenceable how-to  
documentation***

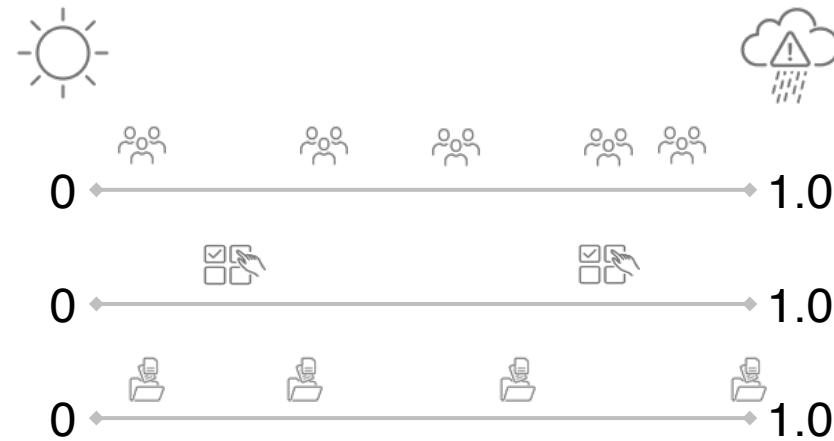
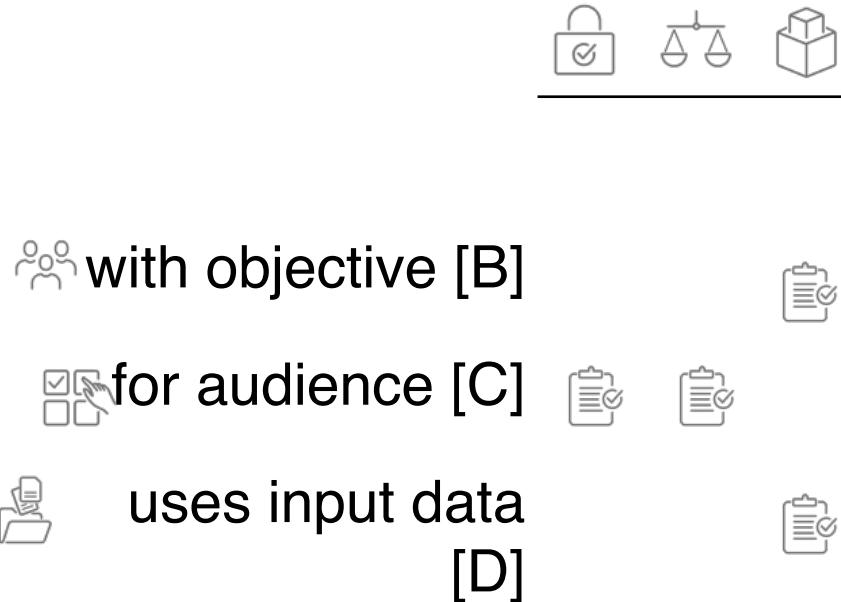


 Privacy

 Legal

 Security

# Risk scoring for reviewers



Objective score = 0.5  
Audience score = 0.8  
Input data score = 1.0

Total risk score = 2.3 out of 3.0

- Privacy (cube icon)
- Legal (scale icon)
- Security (lock icon)

# Transition from chaos

## Slowify

- Stop, review, refine process
- Up-front quality checks
- Track approved tech

## Simplify

- ONLY one entry point
- Use case triage as simple as moving files
- Shortest job next (SJN) scheduling
- Separate tech onboarding
- Optimize known tech over new tech

## Amplify

- Identify high value use cases
- Identify low risk use cases

# Winning and Learning



## YOU SHOULD DO THIS

- ✓ Build early warnings documentation
- ✓ Shift-left traffic light filters
- ✓ All green proxy scenarios
- ✓ Automate where possible
- ✓ [Eat this, not that] tech radar
- ✓ Be transparent



## YOU SHOULD AVOID THIS

- ✓ Don't put more water in the pool than you can drain out.
- ✓ Multiple use case in-take funnels
- ✓ Perfection over MVP

# Help we're looking for

- Developing a Generative AI strategy and how to face the challenges
- Managing enterprise curated model zoos and AI hackathon resources across hundreds of developers
- Learning from AI enterprise platform best practices and productivity tools



Bē

Artwork by Leandro Alzate